

DSAAS

PROBLEM


Alice's structured data sets



Bob's structured data sets

PROBLEM


Alice's structured data sets



Bob's structured data sets

PROBLEM


Alice's structured data sets



Bob's structured data sets

PROBLEM


Alice's structured data sets



Bob's structured data sets

PROBLEM


Alice's structured data sets



Bob's structured data sets

PROBLEM


Alice's structured data sets



Bob's structured data sets

PROBLEM


Alice's structured data sets



Bob's structured data sets

PROBLEM


Alice's structured data sets



Bob's structured data sets

PROBLEM

Alice's structured data sets



Bob's structured data sets

PROBLEM

Collaboration on structured data can be difficult, time-consuming, and frustrating.

DSAAS

A Cloud Service for Persistent Data Structures

P. B. le Roux, S. Kroon and W. Bester

April 6, 2016

Stellenbosch University

OTHER USE CASES

- Scientific Reproducibility

OTHER USE CASES

- Scientific Reproducibility
- Debugging and Teaching Programming

OTHER USE CASES

- Scientific Reproducibility
- Debugging and Teaching Programming
- Session-based Interpreters

BACKGROUND

Ephemeral vs Persistent Data Structures

Ephemeral vs Persistent Data Structures

Types of persistence:

VERSION CONTROL

Ephemeral vs Persistent Data Structures

Types of persistence:

Partial Persistence



VERSION CONTROL


Ephemeral vs Persistent Data Structures

Types of persistence:

Partial Persistence



Full Persistence



VERSION CONTROL


Ephemeral vs Persistent Data Structures

Types of persistence:

Partial Persistence



Full Persistence




Confluent Persistence




PATH-COPYING

Achieve full persistence using a technique called *path-copying*




PATH-COPYING

Achieve full persistence using a technique called *path-copying*




PATH-COPYING

Achieve full persistence using a technique called *path-copying*




PATH-COPYING

Achieve full persistence using a technique called *path-copying*




PATH-COPYING

Achieve full persistence using a technique called *path-copying*



PATH-COPYING

Achieve full persistence using a technique called *path-copying*



HASHED ARRAY MAPPED TRIE (HAMT)



Figure 1: An example of an HAMT. The grey blocks represent the bitmaps, and the white cells represent the array of references to key–value pairs stored in this trie.


HASHED ARRAY MAPPED TRIE (HAMT)



$$h(k_1) = 00000 \dots$$

Figure 1: An example of an HAMT. The grey blocks represent the bitmaps, and the white cells represent the array of references to key–value pairs stored in this trie.


HASHED ARRAY MAPPED TRIE (HAMT)



$$h(k_1) = 00000 \dots$$

Figure 1: An example of an HAMT. The grey blocks represent the bitmaps, and the white cells represent the array of references to key–value pairs stored in this trie.

HASHED ARRAY MAPPED TRIE (HAMT)




$$h(k_1) = 00000 \dots$$

$$h(k_2) = 00011 \ 00011 \dots$$

Figure 1: An example of an HAMT. The grey blocks represent the bitmaps, and the white cells represent the array of references to key–value pairs stored in this trie.

HASHED ARRAY MAPPED TRIE (HAMT)



$$h(k_1) = 00000 \dots$$


$$h(k_2) = 00011 \ 00011 \dots$$

$$h(k_3) = 00011 \ 00101 \dots$$

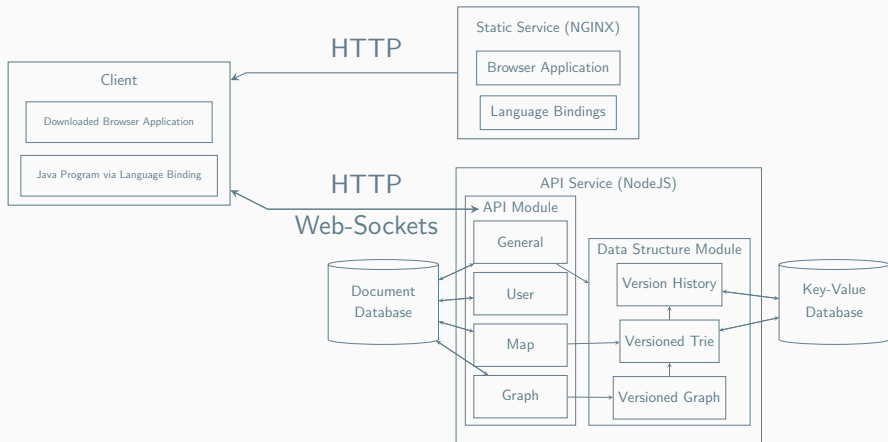
Figure 1: An example of an HAMT. The grey blocks represent the bitmaps, and the white cells represent the array of references to key–value pairs stored in this trie.

DEVELOPMENT

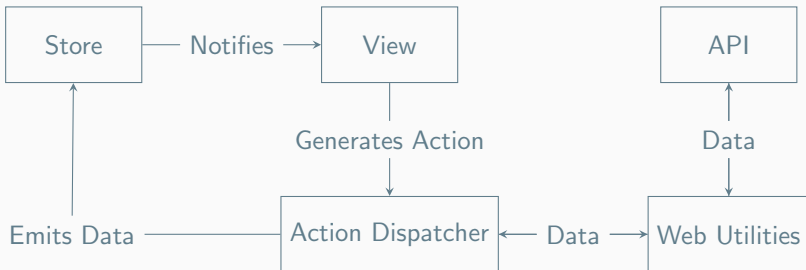
ARCHITECTURE OVERVIEW



ARCHITECTURE OVERVIEW

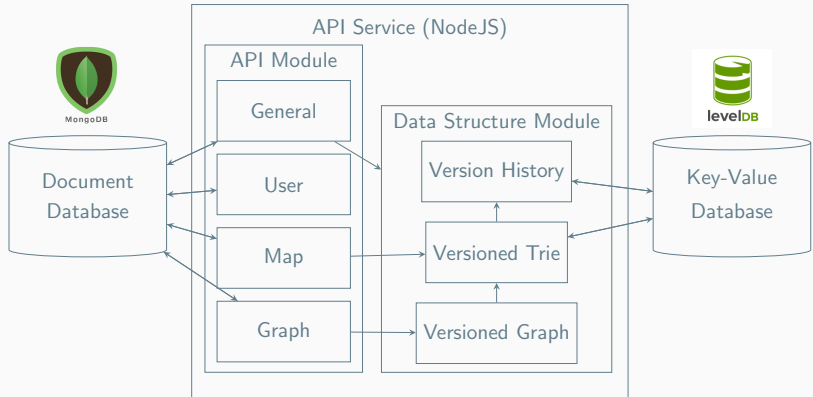


ADAPTED FLUX ARCHITECTURE





BACK END

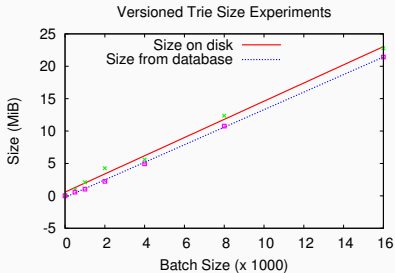


VERSIONED TRIE

- Based on the Hashed Array Mapped Trie (HAMT).
- Implemented on storage instead of in memory.
- Three-Way Merge operation for confluent persistence.
- Detecting transpositions through Zobrist hashing.

EVALUATION

EVALUATION



Insertion: 1 = adding 12
ephemeral data items

Removal: 1 = adding 10
ephemeral data items



Merging: 16 000 x 16 000
elements = increase
of 650 KiB 6000
ephemeral items

LATENCY


Remote Server (Library Binding)	206
Localhost (Library Binding)	7.6
Core (JavaScript)	2

Table 1: The latency (in ms) for the *put* operation using the library binding to connect to a remote server and the localhost, and using JavaScript to test it on the core system.

COMPARED TO DAT



COMPARED TO DAT



CONCLUSION

Questions?

